

# CMASan: Custom Memory Allocator-aware Address Sanitizer

Junwha Hong<sup>1</sup>, Wonil Jang<sup>1</sup>, Mijung Kim<sup>1</sup>, Lei Yu<sup>2</sup>, Yonghwi Kwon<sup>3</sup>, Yuseok Jeon<sup>1</sup>

<sup>1</sup> UNIST, <sup>2</sup> Rensselaer Polytechnic Institute, <sup>3</sup> University of Maryland



Junwha Hong  
qbit@unist.ac.kr

Wonil Jang  
dnjsdlf51@unist.ac.kr

Mijung Kim  
mijungk@unist.ac.kr


Lei Yu  
yul9@rpi.edu

Yonghwi Kwon  
yongkwon@umd.edu

Yuseok Jeon  
ys\_jeon@korea.ac.kr


# Address Sanitizer (ASan)

Address Sanitizer: runtime detection tool for detecting memory bugs with shadow memory



# ASan's Internal Allocator

- ❖ ASan **replaces standard allocators with its internal allocator** to manage objects and redzones



**What happens if the target application uses  
a Custom Memory Allocator that ASan cannot replace?**

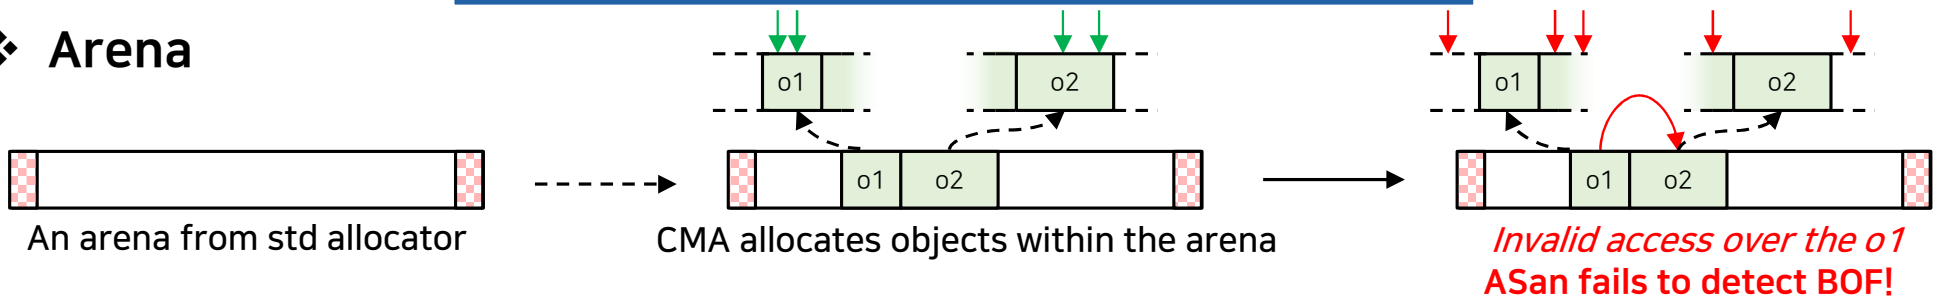
**ASan will miss CMA-allocated objects,  
so bugs affecting them go undetected**

# CMA Patterns Against ASan

- ❖ From 100 C/C++ GitHub projects, identified 78 CMAs in 44 applications
- ❖ Two false negative patterns: Arena (69%) & Recycler (45%)

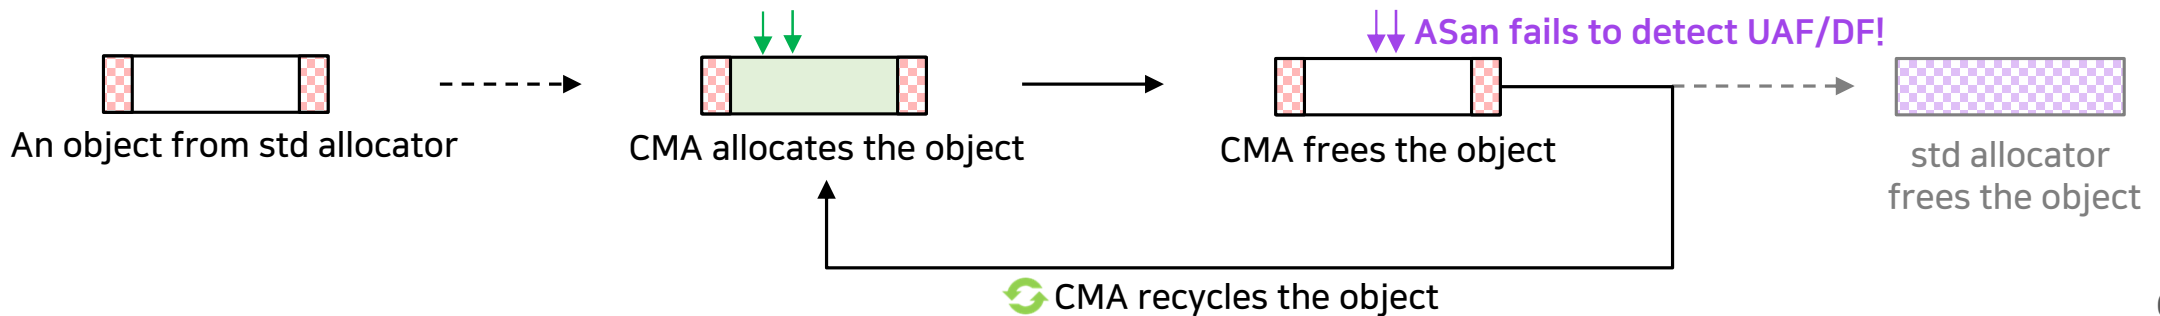
Need **redzones** between the **adjacent objects**!

## ❖ Arena



## ❖ Recycler


Need **Poisoning & Quarantine Zone** before they're recycled!




# Existing Approaches: Shim & ASan API

❖ Shim: provide a mode to switch CMAs to standard allocators

➤ **Not all CMAs are compatible** with standard allocators



*An example of the inplace-realloc API*



*An example of the clear API*


❖ ASan Poisoning API: API for manually poisoning ASan's shadow memory

➤ **Redzone space** must be manually secured


➤ **No Quarantine Zone** support

Both require **manual modification** with a **deep understanding of CMAs**

# CMASan Overview



+ CMA Identification




+ On-demand Metadata Storage


Address Sanitizer




+ CMA API Instrumentation



+ Instance-specific Quarantine zone




+ False Positive Avoidance






# CMASan Overview




+ CMA Identification




+ On-demand Metadata Storage


Address Sanitizer




+ CMA API Instrumentation



+ Instance-specific Quarantine zone




+ False Positive Avoidance




# CMA Identification


- ❖ Extends CodeQL HeuristicAllocationFunction Query → ALLOC candidates  
(size argument flow, additional keywords)
- ❖ Semi-automated Categorization Procedure → collect family functions  
(with helper script for user convenience)



# CMASan Overview




+ CMA Identification




+ On-demand Metadata Storage


Address Sanitizer




+ CMA API Instrumentation



+ Instance-specific Quarantine zone



+ False Positive Avoidance




# CMA API Instrumentation (1/2)


## ❖ ALLOC: secure redzone space

1. Extend **size** before the call

*ptr = cma\_alloc (size);* ***rz\_size***




2. **Poison the right side** of the returned object




3. **Save metadata** into the metadata table using the ***ptr*** as a key

*(ptr, {size, status, cma\_id, ...})* →



## ❖ FREE: retrieve size & poison

1. **Retrieve size** using the argument ***ptr***  
*cma\_free (ptr);*




2. **Check double free** with status

*if status == FREED && cma\_id == (id of cma\_free)* → ***Report Double Free***



3. **Poison and update metadata**




*{size, status, cma\_id, ...}*

# CMA API Instrumentation (2/2)


❖ REALLOC: Compare  $\text{obj}_{\text{old}} \leftrightarrow \text{obj}_{\text{new}}$  to distinguish in-place behavior

1. Extend size before the call

`new_ptr = cma_realloc(old_ptr, size + rz_size);`



2. Poison the right side of the returned object



3. Compare *old\_ptr* and *new\_ptr* and **poison** (*old\_ptr*, *size*) if **different**




❖ CLEAR: record objects & poison them on a CLEAR call

`o1 = cma#1::alloc (size1);`

`o2 = cma#1:: alloc (size2);`


`o3 = cma#1:: alloc (size3);`

`cma#1:: clear();`




Allocation Zone #1~#N  
(push at ALLOC, pop at CLEAR)

# CMASan Overview




+ CMA Identification




+ On-demand Metadata Storage


Address Sanitizer




+ CMA API Instrumentation



+ Instance-specific Quarantine zone




+ False Positive Avoidance




# On-demand Metadata Storage

- ❖ Observation: CMAs **manage their objects** within a relatively **narrow memory region** (e.g., arena)
  - CMASan utilizes a **two-level table** to manage the metadata of CMA objects
  - The 2<sup>nd</sup>-level tables are **allocated on demand** for CMA-managed regions


Requesting metadata for 0x01100010001




CMA-managed region  
(with on-demand 2<sup>nd</sup>-level tables)



# CMASan Overview




+ CMA Identification




+ On-demand Metadata Storage


Address Sanitizer




+ CMA API Instrumentation



+ Instance-specific Quarantine zone



+ False Positive Avoidance






# Instance-specific Free-delaying Quarantine Zone (1/2)

- ❖ Objective 1 (Free-delaying): Delay the recycling of objects while preserving CMAs


*Lifecycle of CMA object*



# Instance-specific Free-delaying Quarantine Zone (1/2)

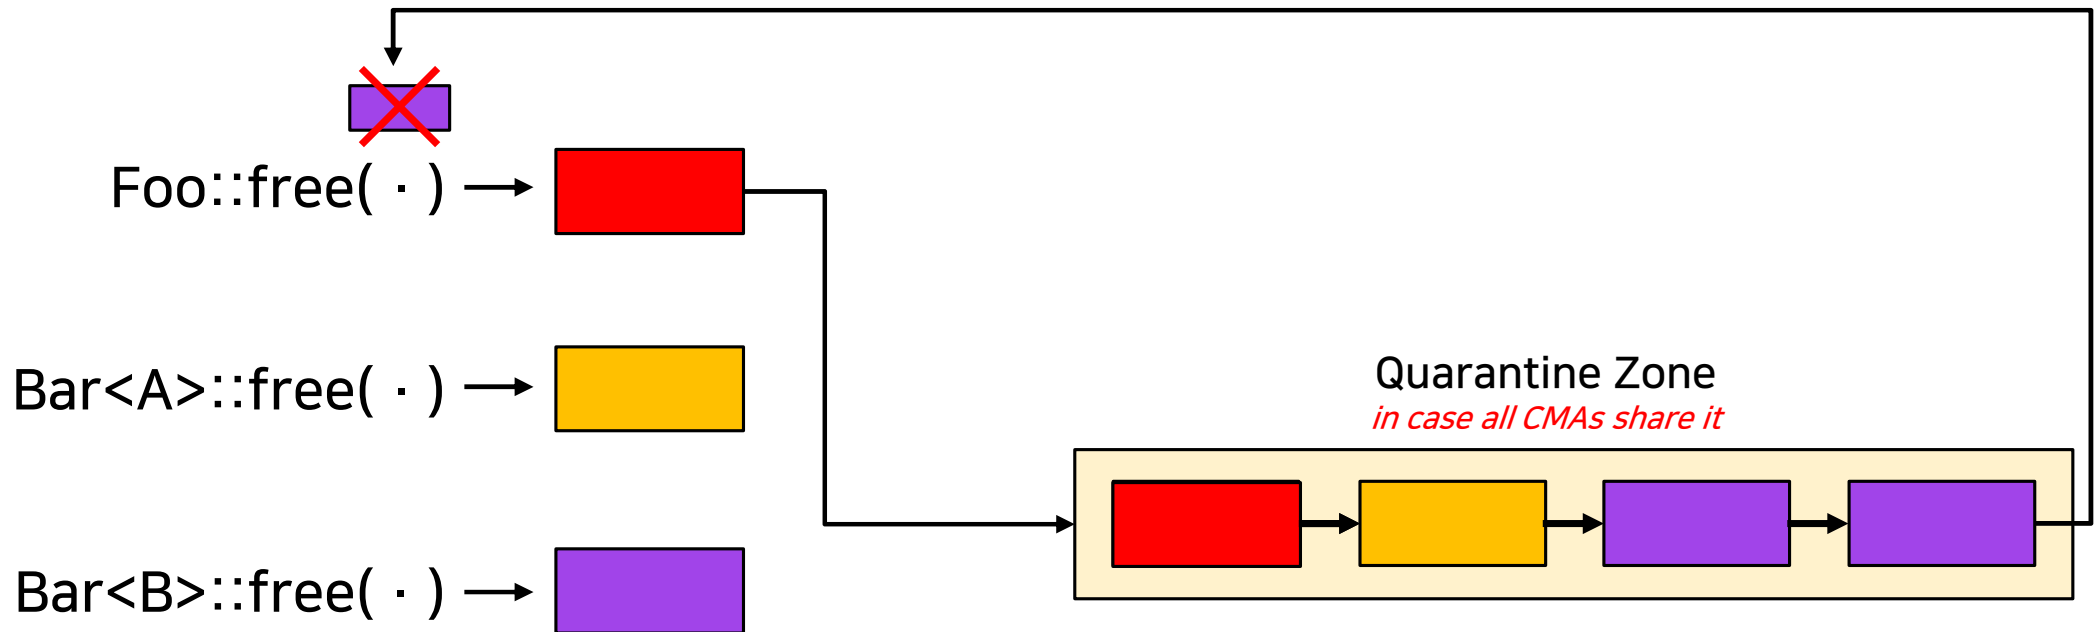
- ❖ Objective 1 (Free-delaying): Prevent the recycling of objects while preserving CMAs

*Lifecycle of CMA object w/ Free-delaying*



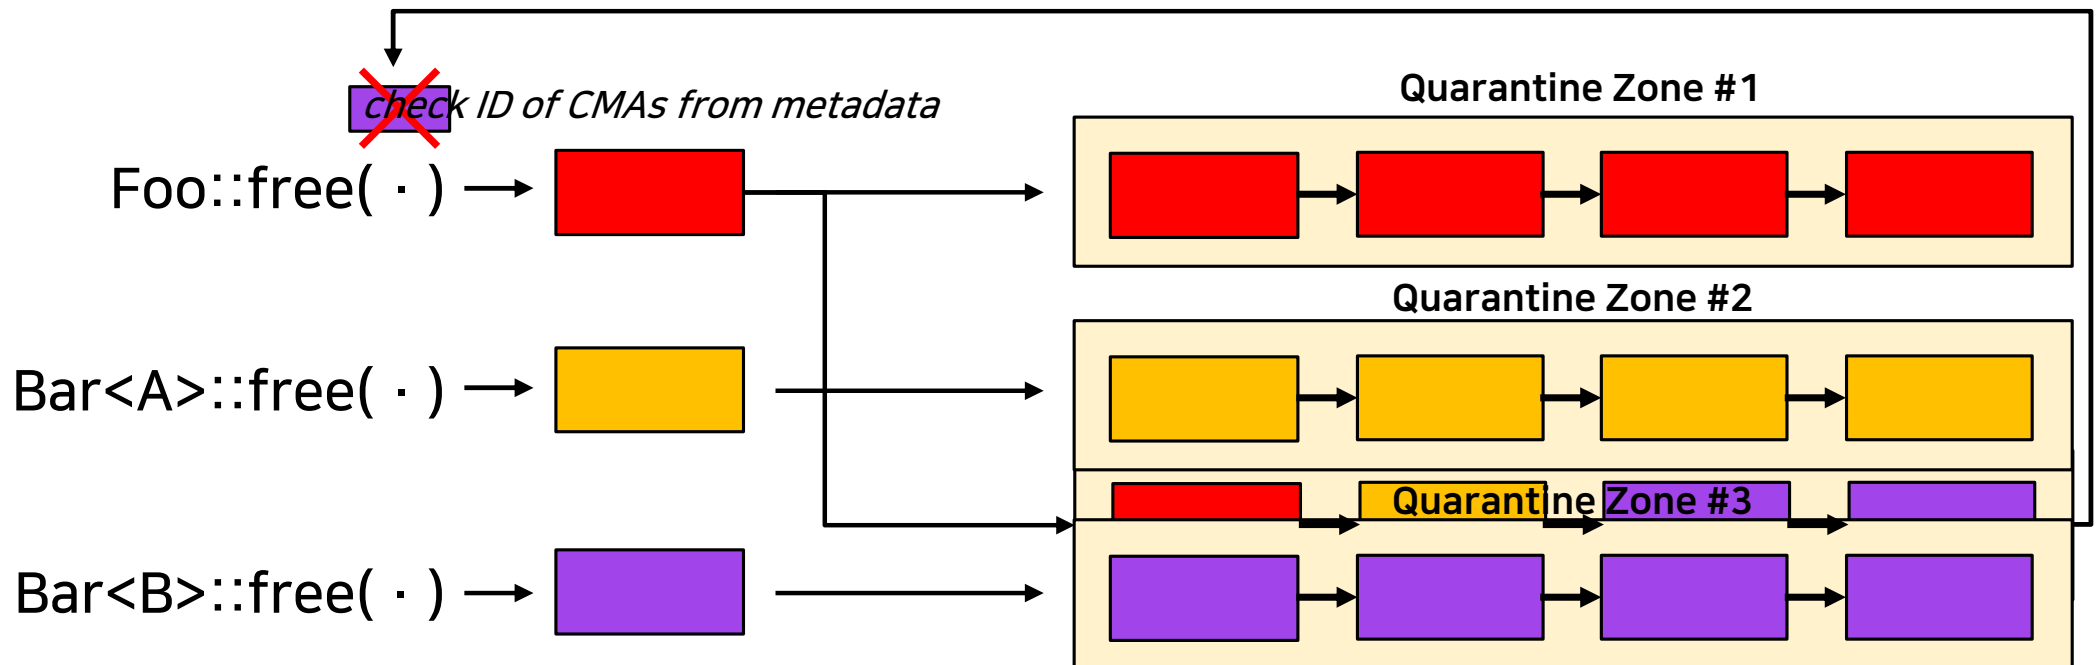
# Instance-specific Free-delaying Quarantine Zone (2/2)

- ❖ Objective 2 (Instance-specific): Distinguish QZ between different CMA instances




# Instance-specific Free-delaying Quarantine Zone (2/2)


- ❖ Objective 2 (Instance-specific): Distinguish QZ between different CMA instances



# CMASan Overview




+ CMA Identification




+ On-demand Metadata Storage


Address Sanitizer




+ CMA API Instrumentation



+ Instance-specific Quarantine zone




+ False Positive Avoidance



# False Positive Avoidance

- ❖ **Suppress reports in CMA APIs** using call stack
  - CMAs often legitimately access poisoned objects (e.g., metadata)
- ❖ Activate only **the outermost CMA's instrumentation** using call stack
  - REALLOC often calls ALLOC and FREE (incorrect redzone sizes/double-free FP)
- ❖ Return the **original object size** on **size-querying API** calls
  - To CMAs, redzones are part of the object



`CMA::get_object_size(obj) - rz_size`


# Evaluation Setup

- ❖ **Target:** Top 12 C/C++ applications that utilize CMAs from GitHub
- ❖ **Comparison Targets:** ASan and Goshawk (Static Analyzer for UAF/DF detection)
- ❖ **Evaluation Metrics:**
  - Detection Coverage
  - Performance Overhead
  - False Positive Avoidance
  - Bug Detection Capability
  - Unknown Bug Detection

# Detection Coverage

Application	CMA Objects
gRPC	623,109,396
PHP	66,431,503
Redis	26,867,307
folly	1,087,644
MicroPython	839,561
ncnn	95,741
cocos2d-x	29,504
Swoole	704
Taichi	468
RocksDB	97,101,151
TensorFlow	92,900
Godot	298

ASan load/store checks on CMA objects




CMA San recognizes many CMA objects overlooked by ASan

Up to 87% of load/store checks are on CMA objects




# Performance Overhead

Performance Overhead (%)



Memory Overhead (%)



CMASan incurs  
1.096x performance and  
1.148x memory overhead  
compared to ASan

# Unknown Bug Detection

Bug ID	Application	Bug Type
<a href="#">CVE-2023-7152</a>	MicroPython	UAF
<a href="#">CVE-2023-7158</a>	MicroPython	BOF
<a href="#">CVE-2023-8946</a>	MicroPython	BOF
<a href="#">CVE-2024-8947</a>	MicroPython	UAF
<a href="#">CVE-2024-8948</a>	MicroPython	BOF
Issue #13004	MicroPython	BOF
Issue #13046	MicroPython	BOF
Issue #13220	MicroPython	BOF
Issue #13428	MicroPython	BOF
Issue #136-1	qhull	UAF
Issue #136-2	qhull	UAF
PR #2213	RapidJSON	BOF
PR #2244	RapidJSON	UAF
PR #2256	RapidJSON	UAF
<a href="#">CVE-2023-7104</a>	SQLite3	BOF
Issue #13230	PHP	UAF
GHSA-rwp7-7vc6-8477	PHP	UAF
Issue #8501	Taichi	BOF
Issue #5734	ncnn	BOF

**CMASan detects  
19 previously unknown bugs**  
including ones undetected for  
**9 years and 2 years** in SQLite3 & PHP  
([6 CVEs](#) , 12 confirmed, 7 patched)

**ASan and Goshawk miss all 19 bugs**  
(no CMA support / missing CMA APIs / complex paths)

# Conclusion

- ❖ CMASan effectively detects all types of CMA-related memory bugs
  - Identifies 19 previously Unknown Bugs
  - 9.63% performance overhead compared to native ASan
  - Extends ASan's coverage without replacing CMAs

# Thank You

[Paper]



CMASan : Custom Memory Allocator-aware Address Sanitizer  
Junwha Hong, Wonil Jang, Mijung Kim, Lei Yu, Yonghwi Kwon, Yuseok Jeon

[Open Source]



CMASan GitHub Repository  
<https://github.com/S2-Lab/CMASan>